Jiya Jaiswal

jaiswaljiya001@gmail.com | Delhi, India | +91 87001 87122 | *LinkedIn*↗ | *GitHub*↗

Education

Kalinga Institute of Industrial Technology (KIIT) — Bhubaneswar, India
B.Tech in Computer Science Engineering, CGPA: 9.26/10 — *Aug 2024 – Oct 2028*

Delhi Public School, Ruby Park — Kolkata, India
AISSCE (Class XII, Science), Percentage: 95% — *Aug 2022 – Jun 2024*

The BSS School — Kolkata, India
WBBSE (Class X), Percentage: 89% — *Jan 2009 – Jun 2022*

Projects

- **OmniDict – Distributed Sharded KV Store | *GitHub*↗** — *Jun 2025 – Aug 2025*
  Go, gRPC, Protocol Buffers, HashiCorp Raft, Consistent Hashing, Docker
  - Achieved **zero data loss under single-node failure** on a live 3-node quorum, verified by a kill-mid-operation test CLI, by implementing **per-shard Raft consensus groups** (HashiCorp Raft) for leader election and log replication.
  - Reduced key remapping to **under 25%** during cluster resizes, measured via rebalance benchmarks, by engineering a **consistent-hash ring with virtual nodes** for even shard distribution.
  - Delivered a **gRPC API** (`Put`, `Get`, `Join`) orchestrated through **Docker Compose** for one-command spin-up of a 3–5 node cluster with reproducible no-loss kill-node tests.
- **MazeMind – Adaptive AI Maze Game | *GitHub*↗** — *May 2026 – Jun 2026*
  C#, Unity, ScriptableObjects, URP, Git
  - Engineered a **real-time adaptation loop** retuning hazards, gem placement, and key spawns per room across **4+ tracked player metrics** (deaths, time-in-section, hazard hits, path choice), by architecting an **AI Director** with priority-sorted `AdaptationRuleSO` ScriptableObjects and a central `AdaptationState` event bus consumed by every room.
  - Delivered **Room 1 end-to-end** (3 sections, spawn points, section triggers, win/lose flow) and a **procedural CheckboxFloorGenerator** producing solvable mazes with guaranteed entry/exit, validated by an in-game `DecisionLogger` capturing every adaptation event for replay and tuning.
  - Shipped **3 fully playable rooms** with escalating mechanics – bullet traps, sliding platforms, forced-fall sequences, NPC enemies, and key pickups – stress-tested across 20+ playthroughs to validate adaptation rule convergence.
- **FactLens – Agentic Misinformation Detector | *GitHub*↗** — *MumbaiHacks 2025*
  Python, FastAPI, Groq LLaMA 3.1, Knowledge Graphs, Docker, GitHub Actions
  - Built a **multi-agent claim verification pipeline** labeling social posts True/False/Partly True at up to **96% confidence**, by chaining two Groq LLaMA 3.1 agents – one for claim extraction, one for evidence-grounded verdict reasoning.
  - Achieved **zero out-of-evidence hallucinations** across all test cases by constraining all LLM reasoning to a **structured knowledge graph** (`kg.json`) and blocking model access to outside context – equivalent to a strict retrieval-augmented generation architecture.
  - Architected the pipeline for **precision/recall/F1 evaluation** across True/False/Partly-True verdict classes; reduced deployment to a single command via **Docker + GitHub Actions CI/CD** with a clean `backend/ extension/ assets/` split for parallel team development.

Skills

Languages: Go, Python, C, C#, Java, JavaScript, TypeScript, SQL
Frameworks & Libraries: FastAPI, gRPC, Protocol Buffers, React, Node.js, Unity, LLM APIs (Groq)
ML & AI: Multi-Agent LLM Pipelines, Retrieval-Augmented Generation, Knowledge Graphs, Model Evaluation (F1/Precision/Recall), Supervised Learning Fundamentals
Systems & Architecture: Distributed Systems, Raft Consensus, Consistent Hashing, Microservices, RESTful APIs
Core CS: Data Structures & Algorithms, Object-Oriented Programming, Operating Systems, Database Design
Dev Tools: Git, Docker, GitHub Actions (CI/CD), Linux, MySQL, Unit Testing

Achievements

- **Top 500 / 3,000+ teams** at MumbaiHacks 2025, Multi-Agent AI track
- **Top 50**, Silicon Labs IoT Challenge 2026 – Smart Healthcare with Edge AI track